

Free Claude Code Plugin Brings Amazon's PR/FAQ Process to Engineers and Founders

punt-labs releases `prfaq`, an open-source tool to generate, review, debate, decide, and iterate on product discovery documents inside the terminal

Press Release

SAN FRANCISCO — March 2027 — Punt Labs released `prfaq`, a free Claude Code plugin that guides engineers and founders through Amazon's Working Backwards PR/FAQ process to evaluate product ideas before building them. The plugin turns a structured discovery conversation into a professional decision-making document — press release, FAQs, risk assessment — then helps stress-test it through automated peer review, simulated review meetings, and directed iteration, all inside the same Claude Code session where the user already works (see [FAQ 15](#)). Engineers who ship products at unprecedented speed can now apply the same rigor to deciding *which* products to ship.

This is a Working Backwards press release describing the intended state as of March 2027. `prfaq` first shipped publicly on 2026-02-16 (v1.0.0) and is at v1.7.3 today; see [FAQ 12](#) for what has actually happened since.

Problem

AI coding tools let a single person build in a weekend what used to take a team a quarter. The bottleneck has shifted: the hard part is no longer building — it is knowing what to build and why (see [FAQ 11](#)). Builders using Claude Code, Cursor, and similar tools ship fast; our working assumption, not yet directly measured, is that many lack formal product training and skip customer definition, competitive analysis, and unit economics — discovering after weeks of building that the product does not resonate.

The consequences are compounding. AI-generated code carries its own quality burden: 66% of developers cite “almost-right” AI output as their top frustration, and 45% report that debugging AI-generated code takes longer than expected [1]. The deeper problem is strategic: our assumption — consistent with, though not directly confirmed by, the “vibe coding” commentary [2] — is that a meaningful share of products built during this wave face rebuilds not because the code is bad, but because the code solves the wrong problem with confidence. Enterprises carry \$1.5–2 trillion in accumulated technical debt [3]; AI-accelerated development risks creating more of the same debt, faster.

A builder who wants product discipline has three options: hire a product manager they cannot afford, read a book and apply frameworks manually, or iterate — ship something, get feedback, ship again. Iteration works. But iteration has real costs: every cycle burns tokens, calendar time marketing to early users, and the founder's energy. Getting customer feedback is often the actual bottleneck. Builders would develop stronger product sense if the frameworks were accessible in the workflow where they already build, not in a book or workshop they will never attend.

Solution

`prfaq` brings product thinking into the builder's environment. Type `/prfaq` in any Claude Code session to start a structured conversation: Claude asks about the target customer, the problem,

the competitive landscape, and the business model. It asks hard questions — the kind a product manager would ask before greenlighting a project.

Claude generates a complete PR/FAQ document: a mock press release describing the product as if it has already shipped, followed by external FAQs (what a customer would ask) and internal FAQs (what leadership, finance, and engineering would ask). The output is a LaTeX-compiled PDF (see [FAQ 6](#)): a clean, professional artifact for sharing with co-founders, investors, or advisors. Not a disposable brainstorm. A decision-making document designed to be read, debated, and revised.

Generation is step one. The plugin then helps the builder stress-test the idea: `/prfaq:review` runs automated peer review against Working Backwards principles and a cognitive bias checklist, surfacing weaknesses the author is too close to see. `/prfaq:research` finds evidence for unsupported claims — scanning local files, indexed documents, and the web — and embeds citations into the document.

Next, debate. `/prfaq:meeting` simulates an Amazon-style PR/FAQ review meeting with agentic personas — a target customer, a principal engineer, a unit economics reviewer, a UX bar raiser, and a skeptical executive — each challenging the document from their perspective. `/prfaq:meeting-hive` takes this further: the personas debate autonomously, surfacing conflicts and consensus without the user moderating every exchange. The user is the PM and final decision-maker; the agents provide cross-functional scrutiny a solo builder otherwise lacks.

Then, decide. `/prfaq:vote` walks the builder through a structured assessment: is the customer problem worth solving, is the solution differentiated, and is now the right time? It surfaces gaps in evidence, identifies cheap ways to reduce uncertainty before committing investment, and renders a verdict. The goal is not a mechanical formula — context always matters — but a thinking exercise that forces the builder to confront what they know, what they assume, and what they have not tested. When the product ships, `/prfaq:externalize` generates a customer-facing press release from the internal PR/FAQ and CHANGELOG.

Finally, iterate. `/prfaq:feedback` takes pointed direction — “focus on CTOs, not solo devs” or “the TAM is too large” — and surgically redrafts affected sections, tracing cascading effects so the user iterates on substance rather than formatting. Together, the commands form a complete product-thinking workflow: generate, stress-test, debate, decide, and iterate.

Customer Quote

“I had three weeks of build planned for a pricing overhaul. The hive meeting killed the feature in twenty minutes — the unit economics reviewer caught that my per-seat pricing broke at the usage tier I was actually targeting. That’s a mistake I would have shipped and found out about from a customer, not from a simulated debate.”

— Priya Raman, Founder, Ledgerly

Getting Started

Install with one command — no account, no subscription, no configuration:

```
curl -fsSL https://raw.githubusercontent.com/punt-labs/prfaq/main/install.sh | bash
```

Then follow the workflow:

Generate. Type `/prfaq`. Claude asks discovery questions about the target customer, problem, differentiation, and business model, then drafts a complete PR/FAQ document. Within an hour, you have a compiled PDF.

Stress-test. Type `/prfaq:review` for automated peer review and `/prfaq:research` to find and cite evidence. Type `/prfaq:meeting` to face a simulated review meeting with agentic personas who challenge the document from every angle — or `/prfaq:meeting-hive` for autonomous multi-agent debate.

Decide. Type `/prfaq:vote` to think through whether the idea is ready to pursue. The command walks through three questions — is the problem worth solving, is the solution differentiated, and is now the right time — surfacing gaps in evidence and judgment along the way.

Iterate. Type `/prfaq:feedback` followed by a directional instruction. The plugin traces cascading effects and surgically redrafts affected sections. Repeat until the document — and the product idea — is ready.

Spokesperson Quote

“The people building products with AI make product decisions every day — they just don’t always have the frameworks. We built `prfaq` because the engineers and founders using Claude Code are also deciding what to build and for whom, and they deserve a tool that meets them where they work. The PR/FAQ process has been proven at Amazon for two decades. We didn’t invent it. We made it accessible inside a terminal.”

— **Jim Freeman**, Founder, Punt Labs

Call to Action

`prfaq` is free and open source, available now at <https://github.com/punt-labs/prfaq>. Install it, type `/prfaq`, and find out if your next product idea is worth building before you build it.

Frequently Asked Questions

External FAQs

Q1: What is prfaq and who is it for?

prfaq is a free Claude Code plugin that guides engineers and founders through the Amazon Working Backwards PR/FAQ process. The primary audience is technical builders who use AI coding tools to create products and want to validate an idea before committing weeks or months of development time. As Claude Code's user base expands beyond professional engineers [4], features like `/prfaq:import` — which enriches an existing document rather than generating one from scratch — extend reach to non-technical users who adopt Claude Code as a thinking partner.

Q2: How is this different from just asking Claude to write me a PR/FAQ?

This is the real alternative, not a template — anyone with Claude Code can type “write me a PR/FAQ for this idea” for free, one prompt, right now. What that single prompt does not do: apply a structured six-phase discovery conversation that surfaces gaps before drafting; run an adversarial peer review pass against a named cognitive-bias checklist after drafting; convene four independently-prompted personas with distinct risk lenses who can disagree with each other in a visible debate, rather than one model generating uniformly agreeable output; trace cascading edits across a dependency graph when you give feedback, so a change to the customer segment does not leave the TAM section stale; or compile to a cited, professionally typeset PDF. A single prompt can approximate the shape of a PR/FAQ. It cannot reproduce a structured process with built-in adversarial review — that requires the orchestration prfaq provides, not better prompting of a single turn.

Q3: How do I get started?

Run the one-line installer, then type `/prfaq` in any Claude Code session:

```
curl -fsSL https://raw.githubusercontent.com/punt-labs/prfaq/main/install.sh | bash
```

The plugin walks you through the process. No account, no API keys, no configuration. Installation takes under a minute.

Q4: What does the output look like?

A professionally typeset PDF compiled from LaTeX. The document includes a press release, external and internal FAQs, and a four-risks assessment. Designed to be shared with co-founders, investors, or advisors — not a throwaway brainstorm.

Q5: Do I need to know LaTeX?

No. The plugin generates the LaTeX source and compiles it automatically. You never see or edit LaTeX unless you choose to. You need a TeX distribution installed (TeX Live or similar); the installer checks for it and provides setup instructions if missing.

Q6: Why LaTeX instead of Markdown?

Three reasons. First, a PR/FAQ is a decision document shared with co-founders, investors, and advisors. LaTeX produces professionally typeset PDFs with consistent typography, precise layout, and visual credibility that Markdown-to-PDF pipelines do not match. The medium signals the thinking is serious. Second, LaTeX's structured environments (`faqpair`, `customerquote`, `spokespersonquote`) give the LLM a precise grammar to write into — the AI reads and writes `.tex`, the human reads the compiled `.pdf`. This separation means the source format can be rich and semantic without burdening the author. Third, LaTeX's package ecosystem (`biblatex` for citations, `mdframed` for callout boxes, `enumitem` for structured lists) provides extensibility Markdown lacks without custom renderers.

The trade-off is friction: LaTeX requires a TeX distribution (~4 GB), and users unfamiliar with it may find compilation errors intimidating. The installer detects TeX absence and provides platform-specific setup instructions, but it does not auto-install the dependency — the user must install it themselves. The plugin handles all LaTeX generation — the user never writes `.tex` directly. This is a two-way door: if user feedback shows the TeX dependency is a meaningful adoption barrier, adding a Markdown output mode is straightforward since content generation is format-independent. The structured thinking matters more than the output format.

Q7: What happens to my data?

Everything runs locally. The plugin is a set of prompt files and a LaTeX template — it does not phone home, collect telemetry, or store data. Your PR/FAQ content is processed by Claude Code using your existing Anthropic API relationship. The generated `.tex` file and compiled PDF stay on your machine. The one exception is the opt-in `/prfaq:feedback-to-us` command, which transmits a 1–5 satisfaction rating and optional comment — no document content, no project data. Data handling is governed by Claude Code's privacy model and your Anthropic terms of service.

Q8: How much does it cost?

Free. `prfaq` is open source under a permissive license. No paid tier, no freemium upsell, no planned monetization. The plugin itself adds zero cost. You need an active Claude Code subscription or API access, which is a separate cost from Anthropic.

For API users, the token cost is small. A full PR/FAQ session — discovery conversation, document generation, and peer review — consumes roughly 100–200K input tokens and 20–40K output tokens across multiple turns (context is re-sent each turn). At Sonnet 5 pricing (\$2/\$10 per million tokens, as of Q3 2026), that is \$0.40–0.80 per initial draft. Each feedback cycle is lighter: roughly 40–80K input tokens and 10–20K output tokens, or \$0.18–0.36 per round. A typical document through three or four revision cycles costs under \$2. Opus sessions cost roughly 2.5× more at list pricing (Claude Opus 5 lists at \$5/\$25 per million tokens, as of Q3 2026).

For comparison, experienced product consultants charge \$100–300/hour. Claude Code subscribers on the Pro plan (\$20/month), Max 5× (\$100/month), or Max 20× (\$200/month) (as of Q3 2026) pay nothing incremental — token usage falls within the subscription's included allowance. Pricing changes; check <https://anthropic.com/pricing> for current rates.

Q9: Can I iterate on the document?

Yes. Type `/prfaq:feedback` with a directional instruction — “reframe the problem around compliance risk” or “the TAM is too small” — and the plugin traces cascading effects and surgically redrafts affected sections. Each feedback cycle recompiles the PDF and runs peer review automatically. You can also edit the `.tex` file directly or re-run the full discovery process. The Working Backwards process is designed for rapid iteration — the point is to iterate on a short document instead of a shipped product.

Internal FAQs

*Stage is **validated** because the product has shipped and six months of real adoption data exist to report — not because the value hypothesis has been validated. It has not; see [FAQ 12](#) and [FAQ 14](#).*

Value & Market

Q10: What is the total addressable market?

Reachable users today. The plugin runs on Claude Code. We do not know how many people use Claude Code. What is public: Claude Code’s annualized run-rate revenue reached approximately \$2.5B as of early 2026, doubling since January [5]. That figure is Claude Code-specific, within Anthropic’s \$14B total ARR. Working from revenue to user count requires one key unknown: revenue per user. A blended ARPU of \$40–100/month across Pro, Max 5×, Max 20×, and API tiers is itself an unvalidated estimate — the true distribution of users across tiers is not public, and API spend per user varies by orders of magnitude. Dividing \$2.5B by an unknown denominator yields a rough order-of-magnitude placeholder, not a derived estimate: somewhere in the range of 2–5M Claude Code users as of Q1 2026. Every digit is soft. Of those, the product-decision-making subset (founders, technical leads, PMs) is an assumed 10–25% — a segmentation ratio we have not validated and plan to test through the validation trial ([FAQ 14](#)). That yields a placeholder range of 200K–1.25M reachable users today.

This is not a TAM-based revenue argument (the plugin is free); it is a reach argument. The question is how many builders the plugin can help, not how many it can monetize. The numbers above exist to establish that the reachable population is large enough to matter, not to claim precision about its size. The trajectory of both Claude Code adoption and the broader vibe-coding wave suggests the range grows significantly over 12–18 months.

Trajectory upside: non-technical builders. A massive wave of non-programmers is building software. Lovable has nearly 8 million users [6]. Replit has 35M+ users and its CEO pivoted away from professional developers [7]. Bolt reached \$40M ARR [8]. Cursor has 1M daily active users [9]. A quarter of startups in Y Combinator’s W25 cohort have codebases almost entirely AI-generated [10]. These builders ship fast; our assumption, not yet directly measured for this segment, is that many lack formal product training and skip customer definition, competitive analysis, and unit economics. The “vibe coding” wave [2] produces unprecedented build velocity with limited product discipline. This is the largest segment and the one with the most acute need for the frameworks [prfaq](#) provides. However, many of these builders work in browser-based tools, not terminals. They are reachable only if Claude Code’s distribution expands to meet them or if they migrate to Claude Code as their projects grow. This segment represents trajectory upside, not launch-day reach.

Expert PMs seeking AI leverage. Experienced product managers and product-minded founders already know frameworks like Working Backwards [11] and four risks [12]. Their constraint is capacity, not knowledge. A PM at a startup wearing four hats does not have time to write a thorough PR/FAQ from scratch. An AI tool that drafts the document from a structured conversation, applies the framework rigorously, and handles formatting lets a single PM do work that previously required a PM plus an analyst plus a designer. This segment is smaller but higher-intent and more likely to produce documents worth sharing.

Beachhead: Amazon alumni and Working Backwards practitioners. Among expert PMs, a distinct micro-segment offers the highest early-adoption conviction: engineers and leaders who practiced Working Backwards at Amazon and now work elsewhere. These practitioners already believe in the PR/FAQ process — they do not need education on the framework or persuasion

of its value. Their gap is operational: they lack the institutional support (templates, trained reviewers, organizational muscle memory) that made the process frictionless inside Amazon. The plugin fills that gap. This micro-segment is small relative to the total addressable audience, but it is the most efficient beachhead: zero education cost, high framework fidelity expectations (which sharpens the product), and natural word-of-mouth to colleagues curious about the process. The author's professional network provides direct access to this group.

The meeting feature changes the value proposition for each segment. The review meeting commands (/prfaq:meeting and /prfaq:meeting-hive) shift what the plugin *is*. For non-technical builders, the meeting is the most transformative feature. Most have never sat in a product review meeting. They have never had a principal engineer question their scaling assumptions, a unit economics reviewer challenge their pricing model, or a UX bar raiser ask how a first-time user would react. The meeting gives them an experience previously available only inside companies with mature product organizations. It is the difference between a tool that helps you write a document and a tool that helps you *think like a product team*. For expert PMs, the meeting is a force multiplier. An experienced PM knows what a review meeting should surface — but running one requires assembling five busy people in a room, often the bottleneck. The autonomous hive variant lets a solo PM get cross-functional scrutiny in minutes, not weeks of calendar negotiation. Across both segments, the meeting is the feature most likely to drive adoption and word-of-mouth: it is visceral, surprising, and difficult to replicate outside a structured multi-agent system.

Reachable versus reached. The 200K–1.25M figure above is a reach estimate, not a forecast, and it has not been reconciled against six months of observed data until now: 24 GitHub stars, 4 forks, and one external pull request as of [FAQ 12](#). Reach was never the binding constraint — the plugin is discoverable and free to install. Conversion is. Whatever the reachable population's true size, the observed conversion rate from it is currently very low, which is a distribution and positioning question, not a market-size question.

Q11: What evidence do we have that customers want this?

Three categories of indirect evidence plus two demand signals from the author's direct network, but no structured primary customer research yet. First, the "vibe coding" phenomenon: thousands of products are built at speed without product discipline; the author has personally observed rebuild discussions in startup communities and forums, though this has not been systematically tallied. Second, the explosion of "how to think about product" content targeting builders — books, courses, and newsletters suggest unmet demand for product skills among people who build with AI tools but lack formal product training. Third, the Working Backwards process has two decades of validation at Amazon [11], where it is mandatory for new product proposals. The gap is not whether the framework works — it is accessibility. Builders do not attend PM workshops. They install plugins.

Fourth, an organic demand signal from the author's professional network: a former colleague who learned the Working Backwards process from the author is introducing it at a new organization — and needs a lightweight way to practice the framework without the institutional support structure Amazon provides. This is the gap the plugin addresses: someone who values the process but lacks organizational scaffolding to run it. The plugin was partly inspired by this pattern: practitioners who leave environments where Working Backwards is institutionalized and want to carry the discipline with them. One anecdote is not validation, but it is a concrete instance of the target user and problem occurring organically.

Fifth, a first-hand adoption signal, as of August 2026 [13]: the author has personally spoken with product managers at two separate companies who report using the plugin, and has received additional unsolicited inbound interest from individuals via LinkedIn. Treated strictly, this is a reach signal, not a value signal — it evidences that the plugin has traveled beyond the author's direct outreach, and nothing more. It does not yet establish which commands these users ran, whether they completed and shared a document, or whether they returned for a second idea — the signals that distinguish adoption from installation (see FAQ 14). Set against the measurable public counters — 24 GitHub stars, 4 forks, and one external pull request (a real installer bug fix, closed without merging) as of FAQ 12 — the reported usage is consistent with a small number of individual practitioners rather than organizational adoption. The next step is to convert these conversations into first-hand written accounts (anonymized) and, ideally, a completed document shared with permission — see FAQ 12 for what this signal does and does not yet support.

Note: we have not validated that technical builders want product frameworks delivered inside Claude Code. Five to ten customer discovery interviews with engineers and technical founders would significantly strengthen this evidence base. The author's professional network of Amazon alumni and Working Backwards practitioners provides a high-conviction recruitment pool — people who already understand and value the framework and can evaluate whether this delivery mechanism adds real value. As the non-technical audience grows [4], a separate round of interviews with that segment would inform whether /prfaq:import resonates.

Q12: What have we learned since launch?

The plugin publicly launched 2026-02-16 (v1.0.0) and is now at v1.7.3, roughly six and a half months and 14 releases later. Development was not continuous: after v1.6.1 (2026-03-20) there was a roughly five-month gap with no release until v1.7.0 (2026-08-13). The key metrics FAQ (FAQ 23) pre-committed to five 6-month targets with an explicit consequence if they were missed. That window closed 2026-08-16. Reporting honestly against it:

1. **GitHub stars:** 24, against a target of 500 (4.8%).
2. **External pull requests:** 1, against a target of 10 — a real installer bug fix from an external contributor, closed without merging.
3. **Installs:** not separately tracked; forks stand at 4, watchers at 1.
4. **Satisfaction ratings via /prfaq:feedback-to-us:** zero submissions to date.
5. **Organic community mentions:** not tracked; none confirmed.

By the letter of the pre-committed consequence, this misses every target and the distribution strategy should be reconsidered. The single "was distribution tried?" question is too coarse, though — the acquisition plan (FAQ 25) named six distinct channels, and they were not all left equally untried. The four that matter most to this question:

- (1) **GitHub discoverability.** Passively executed for 6.5 months by the repo simply being public and topic-tagged. Yield: 24 stars.
- (2) **Amazon alumni network** — the plan's own "highest-conviction," "lowest-cost, highest-conversion" channel. This is the channel that was actually run, deliberately, on warm contacts. Yield: two informal PM conversations, zero completed-and-shared documents of record, zero /prfaq:feedback-to-us submissions, one external pull request (closed without merging).
- (3) **Working Backwards consulting ecosystem.** No record of outreach through this relationship having been initiated.
- (4) **Cold community seeding** (Hacker News, Reddit r/ClaudeAI, Twitter/X). No record of the planned posts having gone out.

Channels 3 and 4 are genuinely untested — silence there says nothing about whether cold and indirect distribution would work. Channel 2, the plan's own highest-conviction bet, *was* executed and converted poorly: warm, pre-sold contacts produced two conversations and no completed documents in 6.5 months. That is a specific, actionable signal, not an absence of data, and it is more concerning than the untested channels because it was the channel expected to work best. It is not enough to run the untested cold channels (3 and 4) — that answers whether public distribution works, but leaves unanswered why the warm channel underperformed. Both questions need an answer: (1) actually execute channels 3 and 4 at the volume originally planned, and (2) find out, from the two known warm contacts, why neither has completed and shared a document — is it the TeX dependency, the time investment, low urgency, or something about the output quality? That diagnostic conversation is cheaper than a cold-channel push and more likely to fix the thing that is actually broken. Value risk is re-rated to High in the interim (FAQ 14); the formal 10–20 person validation trial described there has not been run.

Q13: What is your next step to validate your vision?

Two actions, both dated, now that the original undated 6-month window closed without a check-in (FAQ 12):

- (1) **By 2026-09-30:** hold a direct conversation with each of the two known warm contacts, asking specifically why neither has completed and shared a document — TeX friction, time investment, low urgency, or output quality. Write up both conversations, anonymized, in research/. If neither conversation happens by this date, that itself is the finding: the warm channel is not just underperforming, it is not being worked.
- (2) **By 2026-10-15:** post to at least one of Hacker News, Reddit r/ClaudeAI, or Twitter/X — the cold channels named in FAQ 25 that have no record of ever having been tried — and report the resulting installs, stars, and comments, however small.

If (1) surfaces a fixable friction point, fix it before (2), so the cold-channel post is not sending traffic at a conversion problem that is already diagnosed and unaddressed. If neither action produces a usable signal by 2026-10-15, escalate to the branch-(c) kill-signal logic in FAQ 14.

Q14: How will we validate that builders actually want this?

The customer evidence FAQ (FAQ 11) is honest: we have strong indirect evidence but no primary customer data. The validation plan is straightforward.

Phase 1: Hands-on trials (weeks 1–4). Ask 10–20 target users to install the plugin and use it on a real product idea. Recruit deliberately across the spectrum: (1) casual builders and vibe coders from the punt-labs network and non-technical Claude Code communities, (2) engineers and technical founders from Claude Code channels (Discord, Reddit r/ClaudeAI) and Hacker News “Show HN” threads, (3) experienced product managers or product-minded founders who use structured frameworks, and (4) Amazon alumni and Working Backwards practitioners from the author's professional network — the highest-conviction recruits because they already value the framework and can evaluate whether the plugin faithfully delivers it. This fourth group is uniquely valuable for early feedback: they will judge the plugin against their lived experience with the process, not an abstract expectation. This segmented recruitment is essential to test for the positioning gap — whether the plugin is too rigorous for casual builders and too untrusted for professionals — while the alumni cohort provides a baseline of users who need no persuasion on the framework's value. The ask: try /prfaq on your next idea, run the full workflow through /prfaq:meeting or /prfaq:meeting-hive, and report what happened. The trial must cover the meeting feature explicitly — it is the primary differentiator and the feature with the most

uncertainty about whether users find simulated cross-functional debate genuinely useful versus theatrical.

Phase 2: Qualitative feedback (concurrent with trials). The plugin ships a feedback command (`/prfaq:feedback-to-us`) — a 1–5 satisfaction rating with an optional comment. This is the first feedback signal. During the trial, we ask participants to rate after each completed document. No document content is transmitted; the plugin’s default is zero telemetry.

Phase 3: Structured feedback channel (exploratory). If the basic 1–5 rating proves too shallow, we will expand the feedback command to capture structured responses — what worked, what did not, and whether the output was useful enough to share — via a single API call. This reduces friction of collecting qualitative data from terminal-native users unlikely to fill out a Google Form. The command would remain opt-in, clearly labeled, and would transmit only the user’s explicit responses. This is exploratory; we build it only if Phase 1 reveals users want to give richer feedback than a 1–5 rating.

Signals we are looking for:

1. **Completion rate** — do users finish a full PR/FAQ, or abandon mid-session? Target: 60%+ completion among trial participants.
2. **Shareability** — do users share the generated PDF with a co-founder, investor, or advisor? Target: 30%+ of completed documents are shared.
3. **Decision impact** — did the PR/FAQ change the user’s decision about what to build? Even “I decided not to build it” is a strong positive signal.
4. **Repeat usage** — do users run `/prfaq` on a second idea? Repeat usage within 30 days is the strongest signal of product-market fit for a free tool.
5. **Unprompted referral** — do trial users recommend the plugin to others without being asked?
6. **Meeting utility** — do users who run `/prfaq:meeting` or `/prfaq:meeting-hive` report that the simulated debate surfaced a blind spot they had not considered? Do they change the document as a result? Target: 50%+ of meeting users identify at least one substantive insight they would not have found alone. This is the critical signal for the meeting feature’s value proposition.
7. **Positioning gap** — do casual builders (vibe coders, non-technical founders) abandon the process because it feels too rigorous, while experienced PMs dismiss the output because they do not trust AI-generated strategic analysis? The trial must recruit from both ends of the spectrum, not just the technical middle. If both segments reject the plugin for opposite reasons — too heavyweight for one, too untrusted for the other — the product has a positioning problem that no amount of iteration can fix without a fundamental scope or audience decision.

Pre-committed responses by outcome. The positioning gap (Signal 7) has three distinct failure modes, each requiring a different response. We commit to these branches now — before data arrives — to prevent post-hoc rationalization.

- (a) **Casual builders reject, alumni and PMs engage.** If vibe coders and non-technical founders abandon the process as too heavyweight, but Amazon alumni and experienced PMs complete documents, share them, and return for a second idea — narrow the product to the expert segment. Drop the casual-builder positioning entirely. The plugin becomes a force multiplier for practitioners who already value structured product thinking, not an educational tool for builders who do not. This means the TAM contracts to the expert PM and alumni segments described in [FAQ 10](#), and the non-technical trajectory upside is deferred indefinitely.

- (b) **PMs distrust the output, casual builders engage.** If experienced product managers dismiss the AI-generated analysis as unreliable for real decisions, but casual builders complete documents and find the process valuable — the problem is trust calibration, not product-market fit. Pivot development priority to explainability and framework fidelity: show users *why* the plugin made each recommendation, surface the Working Backwards principles behind each section, and make the framework's logic transparent rather than opaque. The goal is to earn expert trust by demonstrating that the output is grounded in a proven methodology, not generated from generic AI reasoning.
- (c) **Both segments reject simultaneously.** If casual builders find the process too rigorous *and* experienced PMs do not trust the output — dual rejection is a kill signal. The product occupies a positioning no-man's-land that cannot be resolved by iteration within the current delivery mechanism. Response: discontinue active development, or pivot to a fundamentally different delivery mechanism (e.g., a lightweight pre-build checklist instead of a full PR/FAQ, or an integration into a browser-based tool where the non-technical audience already works). Do not attempt to split the difference by making the process simultaneously lighter and more trustworthy — that path produces a product that is mediocre for both audiences.

The choice among these branches depends on which signal fires first and how strongly. A weak signal from one segment with silence from the other warrants patience; a strong signal from both segments warrants immediate action. The alumni cohort described in [FAQ 10](#) is the leading indicator: their response arrives first (warmest leads, highest intent) and calibrates expectations for the broader segments.

Timeline: Recruit the first 10 trial users within 2 weeks of public launch. Complete the 10–20 person trial within 6 weeks. Synthesize findings into a revised customer evidence base. If fewer than 40% of trial users complete a PR/FAQ, or if qualitative feedback consistently indicates the output is not useful enough to share, upgrade the value risk rating to High and reconsider the product strategy.

Status, six months post-launch: this structured trial was not run against the timeline above, which was anchored to a launch date that has now passed and cannot be met. The value risk rating has been upgraded to High in the Risk Assessment regardless — not because the trial ran and failed, but because an unvalidated value hypothesis combined with a missed adoption window ([FAQ 12](#)) is itself a High-risk position under this plan's own logic. Rather than restart this trial on a new deadline immediately, the two dated diagnostic actions in [FAQ 13](#) (2026-09-30 and 2026-10-15) come first — if they surface a clear friction point, fixing it may be cheaper and faster than recruiting 10–20 new trial users; if they do not, this structured trial is the fallback and should be scheduled with a new dated timeline at that point.

Q15: Who are the competitors and why will we win?

The most credible alternative is asking Claude directly to write a PR/FAQ, no plugin at all ([FAQ 2](#)). What that single prompt cannot replicate: a multi-persona review meeting where four experts with distinct risk lenses debate your document and surface blind spots. The meeting commands (`/prfaq:meeting` and `/prfaq:meeting-hive`) are the widest moat — they require orchestrated multi-agent architecture, calibrated persona prompts, and a structured debate protocol a single-turn chat cannot approximate. Equally difficult to replicate: the `/prfaq:vote` command's structured thinking exercise, which walks through evidence gaps, surfaces cheap ways to reduce uncertainty, and forces the builder to confront assumptions — a calibrated decision discipline, not a one-shot prompt. Beyond the meeting and vote, a single prompt lacks:

structured reference guides that catch common mistakes, the four risks evaluation framework, automated peer review against cognitive biases, LaTeX compilation to a shareable PDF, and a repeatable process that improves across uses. The plugin encodes twenty years of Working Backwards practice into a reusable workflow; a single prompt starts from zero every time.

Traditional competitors include product management tools (Productboard, Aha!, Notion templates) and business planning tools (Lean Canvas generators). None are integrated into the builder's working environment. They require context-switching: leave the current tool, open a browser, fill in a form. A Claude Code plugin runs in the same session where the builder already works — the shift from “building” to “product thinking” happens in a single command, not a tab switch. None offer a simulated review meeting.

The most credible competitive response is Anthropic itself. First, Anthropic could feature the plugin — promote it in their plugin directory, recommend it in onboarding, or highlight it as a case study. If that happens, we get distribution we cannot buy, with no downside. Second, Anthropic could build a native product-thinking feature into Claude Code. This scenario is mixed, not simply positive: it would validate the category — the idea that terminal-based AI tools should help users decide *what* to build, not just *how* to build it — but it would also, per the Viability row below, eliminate this product's distribution channel. Category validation is not the same outcome as this product surviving; only the first scenario is good for `prfaq` specifically. The strategic exposure is real — Anthropic controls the plugin system, the distribution, and the competitive landscape.

Market update, Q3 2026. Anthropic launched an official Cowork/Claude Code plugin marketplace in January 2026 (11 plugins, expanded to 24+ by February 2026), including a product-management plugin that generates PRDs, roadmaps, and stakeholder communications and analyzes competitors and metrics [14]. It does not implement the Working Backwards PR/FAQ format specifically, nor a multi-agent debate feature comparable to `/prfaq:meeting//prfaq:meeting-hive`. Separately, a third-party open-source plugin ships a `/pmprompt:working-backwards` command that produces a PR/FAQ-style spec as one of 26+ bundled PM frameworks [15] — direct overlap on the core framework, from outside Anthropic.

“No one else does this” is not true of the Working Backwards framework itself — `pmprompt` implements it too. The combination not found elsewhere as of Q3 2026 is the framework paired with a multi-agent review-meeting feature and an automated peer-review/citation pipeline.

The plugin encodes Amazon's Working Backwards process specifically [11], with reference guides, a peer review framework grounded in Cagan's four risks [12], and a Kahneman-informed decision quality checklist — but this is not the only implementation of that framework, and a general-purpose PM plugin from Anthropic itself now exists. Whether users who value the specific framework and the meeting feature would continue to prefer this tool over Anthropic's own official plugin or the `pmprompt` alternative is unmeasured. The mitigation against platform concentration is explored in the Feature Appendix: alternative CLI-based distribution channels reduce single-platform dependency over time.

Technical

Q16: What are the major technical risks?

The full feature set through v1.7.3 is shipped, including multi-agent review meetings (v0.6.0–v0.8.0), go/no-go investment verdicts with evidence decomposition and investment proportionality calibration (v1.0.0–v1.1.0), and customer-facing press release generation (v1.0.0). The plugin is a set of markdown files plus a LaTeX template and a shell script. No backend, no API,

no database. The LaTeX dependency requires a TeX distribution installed (~4 GB); the installer detects TeX absence and provides platform-specific setup instructions but does not auto-install the dependency.

Multi-agent orchestration — the primary technical challenge — is solved and shipping. The remaining refinement risk is persona fidelity across Claude model updates: each agent must stay in character, disagree substantively (not performatively), and produce debate that is genuinely useful rather than theatrical. If a model update degrades persona quality, the meeting commands could regress from “simulated review meeting” to “five versions of the same generic feedback.” Mitigation: prompt regression testing after model updates and persona calibration from user feedback.

The remaining risk from platform dependency — Anthropic controls the plugin system, distribution, and competitive response — is strategic, not technical. It affects the business’s ability to reach users, not the team’s ability to build the product. That concern is addressed under viability risk, not here.

Q17: What dependencies exist on other teams or systems?

The plugin depends on Anthropic’s Claude Code plugin system for loading and execution, and on a local TeX distribution for PDF compilation. Both are external dependencies outside our control. The Claude Code plugin API is stable (v1), and TeX Live is a mature, 30-year-old ecosystem. If Anthropic changes the plugin format, migration is straightforward — the plugin is a handful of files with simple structure. If TeX Live is unavailable, the plugin still generates the `.tex` source file; only PDF compilation is affected.

All output is portable. The `.tex` source files and compiled PDFs are standalone artifacts — they do not require the plugin to read, edit, or compile. If the plugin becomes unavailable, users retain their documents and can compile them with any standard TeX distribution. No content is locked in a proprietary format or dependent on a running service. This portability is deliberate: the plugin is a workflow tool, not a platform, and the documents it produces belong to the user.

Q18: If this succeeds, what breaks first?

Nothing at the infrastructure level — the plugin is stateless, runs locally, and requires no servers. At scale, the pressures are human, not technical: (1) GitHub issue volume — at 1,000+ users, support requests and feature demands will exceed solo maintainer capacity; mitigation is clear contribution guidelines and aggressive scope discipline (see Won’t Do list); (2) prompt maintenance — as Claude models evolve, skill prompts may need tuning to maintain output quality; (3) backwards compatibility — if Anthropic changes the plugin system, the plugin must migrate; the simple file-based architecture makes this low-risk; (4) community expectations — a popular free tool attracts requests for features that contradict the design philosophy (dashboards, collaboration, web interfaces). The Won’t Do list makes these decisions in advance, not under pressure.

Q19: What is the estimated development timeline?

The full feature set through v1.7.3 is built and shipping. This includes `/prfaq:meeting`, which shipped in v0.6.0 with four agentic personas and visible debate threading, and its autonomous variant `/prfaq:meeting-hive`, which shipped in v0.8.0 using multi-agent orchestration (migrated to Claude Code Agent Teams in v1.2.0). `/prfaq:vote` shipped in v1.0.0 as a three-gate go/no-go decision framework; v1.1.0 added evidence decomposition (three layers per gate), a free evidence pre-check, and Gate 3 investment proportionality testing —

calibrations sourced from running vote on a real investment decision. `/prfaq:externalize` shipped in v1.0.0 for customer-facing press release generation. `/prfaq:import` shipped in v1.0.0 for converting an existing document into the full workflow. `/prfaq:export` shipped in v1.5.0, adding a pandoc-based Word (.docx) output path. The core workflow (`/prfaq`, `/prfaq:feedback`, `/prfaq:review`, `/prfaq:research`, `/prfaq:feedback-to-us`) and supporting infrastructure (stage awareness, version tracking, cross-references, reference guides) are shipped. Remaining planned features (comparative analysis, template customization, alternative CLI distribution) are lower priority and carry no technical risk. No infrastructure to build — remaining complexity is in prompt engineering and skill orchestration. Refinement risks are limited to prompt tuning as Claude models evolve and maintaining persona fidelity across model updates in the meeting and vote commands.

Business

Q20: What is the revenue model?

None today, by design. `prfaq` is free and open source under a permissive license. The strategic value is threefold: (1) it establishes `punt-labs` as the reference implementation for AI-assisted product development tooling — the plugin that defines how structured product thinking works inside an AI coding agent; (2) it builds reputation and visibility in the Claude Code ecosystem; (3) it creates a feedback loop for understanding how builders approach product thinking, which informs other products at `punt-labs`. If validation data eventually warrants it, a conditional commercial path exists (see [FAQ 21](#)).

Q21: Is there a commercial model?

Not today, and not without evidence. The plugin is free and will remain free. However, if product-market fit is confirmed through the validation plan (see [FAQ 14](#)), there is a natural premium tier: subscription access to professional research sources. The free plugin uses web search to find evidence for claims. A premium research tier could provide access to professional databases — Pitchbook for market sizing, Statista for industry statistics, proprietary industry reports — as well as AI-powered survey agents that conduct lightweight customer discovery interviews on the user's behalf. These capabilities address the weakest link in most PR/FAQs: the evidence base. A builder who can generate a PR/FAQ with real market data and structured customer signals has a meaningfully better decision document than one relying on free web search.

This is a path, not a plan. The decision to build a commercial tier is gated on three conditions: (1) the free plugin achieves PMF — repeat usage, shareability, and community traction as defined in [FAQ 14](#); (2) user feedback identifies research quality as a bottleneck; (3) the unit economics of licensing professional data sources close at a price point the target audience will pay. Until all three conditions are met, the commercial model remains a documented option, not a roadmap item.

Q22: What does the cost structure look like at steady state?

`prfaq` has zero infrastructure cost — no servers, no databases, no APIs. The real cost is maintainer time. At current scale (six and a half months into public availability, well below the 500-star target, with a roughly five-month development gap in the middle of that window — see [FAQ 12](#)): well under 2 hours per week on average, though unevenly distributed rather than steady. At

the 500-star target: 4–8 hours per week covering issue triage, community engagement, prompt maintenance as Claude models evolve, and occasional LaTeX template updates. At 1,000+ users: maintainer capacity becomes the constraint. The plugin is designed to be self-sustaining: the entire codebase is markdown files, a LaTeX template, and a shell script — no servers, no API keys, no databases, no backend. Install once, use forever. If punt-labs disappears tomorrow, the plugin continues working. This zero-infrastructure architecture makes it trivially forkable. If maintainer capacity is exceeded, the community can maintain it independently. At that scale, the project either recruits community co-maintainers or narrows support scope to bug fixes only. The opportunity cost is hours not spent on other punt-labs projects. This is acceptable while the plugin establishes punt-labs as the reference implementation for AI-assisted product development tooling; re-evaluate if other punt-labs priorities require the same hours.

Q23: What are the key metrics for success?

Five metrics, measured via GitHub analytics, installer telemetry, plugin feedback, and community signals:

1. **Install count** — target: 500 installs within 6 months. Intended measurement: GitHub release download counts and install script fetch logs; neither has been read against this target (see [FAQ 12](#)).
2. **User satisfaction ratings** — target: average 4.0+ out of 5.0 from users who run the feedback command. Shipped in v0.7.0; captures a 1–5 score with an optional comment.
3. **GitHub stars and forks** — target: 500 stars within 6 months. Measures discoverability and perceived value.
4. **Issues and pull requests from external contributors** — target: 10 external PRs within 6 months. Measures whether the plugin is useful enough that people invest time improving it.
5. **Mentions in builder and developer communities** (Hacker News, Twitter/X, Reddit, Discord servers) — target: 3 organic mentions within 3 months. Measures whether the plugin solves a real problem people tell others about.

If none of these targets are met after 6 months, the plugin is not reaching its audience and the distribution strategy should be reconsidered. This window closed 2026-08-16 — see [FAQ 12](#) for the actual results and what they mean.

Q24: Why now? What has changed?

Three shifts converged. First, AI coding tools reached mainstream adoption in 2025–2026, creating a large population of engineers and technical founders who can build products independently but lack product discipline. Claude Code alone accounts for 4% of GitHub public commits, with projections of 20%+ by end of 2026 [16] — the volume of AI-assisted building is accelerating. Second, Claude Code’s plugin system matured, making it possible to deliver structured workflows inside the developer’s existing tool. Third, the consequences of “vibe coding without product thinking” are becoming visible — failed launches, rebuilds, wasted months — creating demand for lightweight product frameworks that do not require hiring a PM or leaving the terminal. A fourth tailwind is emerging: Claude Code is crossing over to non-technical users [4], expanding the potential audience beyond the technical core.

Q25: What is the customer acquisition strategy?

Distribution is the product strategy for a free plugin. The original plan named six channels. No paid acquisition. Six and a half months post-launch, here is each channel's actual execution status, not the plan for it:

- (1) **GitHub discoverability** (README, topics, stars). Passively executed simply by being public — no active promotion of the listing itself. Yield: 24 stars, 4 forks.
- (2) **Amazon alumni network** — the plan's highest-conviction channel: warm contacts who already believe in Working Backwards and need only be sold on the delivery mechanism. This is the one channel where deliberate outreach happened. Yield: two informal conversations (FAQ 11), zero completed-and-shared documents, zero /prfaq:feedback-to-us submissions, one external pull request closed without merging.
- (3) **Working Backwards consulting ecosystem** (indirect, via the author's relationship with the methodology's author). No record of outreach through this relationship having been initiated.
- (4) **Cold community seeding** (Hacker News, Reddit r/ClaudeAI, Twitter/X). Not executed — no record of any planned post having gone out.
- (5) **Content marketing** (this PR/FAQ itself as a case study and demo artifact). The artifact exists and is public in the repository, but has not been actively shared or submitted anywhere as content marketing.
- (6) **Word-of-mouth** (shared PDF output linking back to the repository). Not measurable with current instrumentation — no referral tracking exists to distinguish a word-of-mouth install from any other.

The secondary audience — non-technical Claude Code users [4] — was not actively targeted in V1, by design; this is spillover, not acquisition strategy, and remains true.

The estimated time investment of 2–4 hours per week on community engagement for the first 3 months was not sustained — see the roughly five-month development gap noted in FAQ 12. The Hacker News conversion assumption (a front-page post reaching 20–30K views, converting at 0.1–0.2%, yielding 20–60 installs) remains an untested assumption, not a forecast, since no such post has been made; FAQ 13 commits to testing it by 2026-10-15. The original “if organic channels do not work within 6 months, revisit the distribution hypothesis” pre-commitment is superseded by the per-channel accounting above and the dated actions in FAQ 13 — that window has closed and is not being reopened undated a second time.

Q26: What are we not building?

Not a SaaS product management platform. Not a competitor to Jira, Linear, Productboard, or Notion. No collaboration features, dashboards, or analytics. No web interface. The plugin is a single-player tool that runs locally, produces a document, and gets out of the way. Scope discipline is the feature. See the Feature Appendix for the full Shipped / Planned / Won't Do breakdown.

Q27: It is one year from now and prfaq failed. What went wrong?

Most likely failure scenario: builders do not want product discipline, even when free and frictionless. The vibe coding audience builds for the joy of building, and a tool that asks “what evidence do you have that customers want this?” feels like a buzzkill, not a feature. Second scenario: the plugin does not surface in a crowded marketplace — the Claude Code plugin ecosystem is growing rapidly, and discoverability is poor. If users cannot find prfaq among

hundreds of alternatives, organic adoption stalls regardless of quality. Third scenario: the output quality is not high enough to share with stakeholders — if the generated PR/FAQ reads like AI slop rather than a credible decision document, users will not trust it for real decisions. Fourth scenario: the PR/FAQ process is too heavyweight for the target audience — builders want a quick sanity check, not a formal document, and a lighter-weight alternative captures the market. Fifth scenario: maintenance burden — the shipped meeting commands (`/prfaq:meeting` and `/prfaq:meeting-hive`) require ongoing persona calibration across Claude model updates, consuming maintainer time that should go toward community engagement and adoption. Sixth scenario: we misidentified the customer or their motivation — the tool solves the wrong person's problem, or targets the right person for the wrong reason. Builders may want stakeholder theater (a polished-looking document that signals rigor) rather than genuine product discipline, and when `prfaq` forces real answers to hard questions, they abandon it for something that produces impressive artifacts without uncomfortable introspection.

The scenario that cannot be iterated away. Scenarios 1–6 are product problems: each can be diagnosed and addressed by changing the plugin — improving output quality, simplifying the process, fixing discoverability, recalibrating persona fidelity, sharpening the customer definition. The seventh scenario is different in kind. Unlike the other scenarios, this failure mode cannot be fixed by trying harder.

Seventh scenario: the product falls between two stools — too rigorous for vibe coders, too untrusted for professionals. Casual builders who thrive on intuition and speed find the structured PR/FAQ process heavyweight and prescriptive: it demands customer definitions, competitive analysis, and unit economics from people who just want to ship. Meanwhile, experienced product managers and executives do not trust AI-generated strategic analysis for real decisions — they view the output as a parlor trick, not a credible basis for resource allocation. The result is a positioning gap where neither audience fully adopts the tool: the casual segment bounces off the rigor, the professional segment bounces off the provenance. This is not a deficiency in any single feature or flow. It is a market-positioning problem: the plugin is aimed at a segment that may not exist in sufficient density — builders who want *exactly this level* of rigor, delivered *exactly this way*. No amount of product iteration resolves a gap in the audience itself. Resolution requires either repositioning (choose one end of the spectrum and serve it fully) or narrowing to a single segment where the rigor-to-trust ratio is already right.

The Amazon alumni beachhead described in [FAQ 10](#) is the direct test of whether that segment exists. These practitioners already believe in the framework and already trust structured product discipline — they do not bounce off the rigor. If even this highest-conviction cohort does not adopt the plugin, the positioning gap is confirmed and no broader audience will either. Conversely, if the alumni cohort adopts enthusiastically, they define the “right-rigor” segment and the product can expand outward from that base rather than straddling two audiences that want fundamentally different things.

Risk Assessment

Risk	Rating	Assessment
Value	High	Re-rated from Medium: the 10–20 person validation trial committed to in FAQ 14 was not run in the six months since launch, and the pre-committed 6-month metrics window closed with targets missed by an order of magnitude (FAQ 12). Market-level evidence remains strong (proven framework, large technical audience). Customer-level evidence is thin but no longer zero: two first-hand PM conversations plus inbound interest (FAQ 11), not yet a structured trial. Positioning risk persists — the plugin may be too rigorous for casual builders and too untrusted for professional PMs, falling between two stools. Plugin is free, so adoption barrier is low. Mitigation: the two dated actions in FAQ 13 (diagnose the underperforming warm channel by 2026-09-30, test a cold channel by 2026-10-15), and the structured trial from FAQ 14 if those do not resolve the open question.
Usability	Medium	One-line install, one slash command, Claude guides interactively. TeX dependency is the main friction: the installer detects TeX absence and provides platform-specific setup instructions but does not auto-install the ~4GB dependency — the user must install it themselves. Non-technical users (a claimed growth segment in FAQ 10) have not been tested on the full onboarding path from zero to compiled PDF. Time to value is designed for under one hour but is unmeasured. Rating will be revised to Low when the validation trial confirms median time-to-first-PDF under one hour across environment types (macOS, Linux, WSL).
Feasibility	Low	Full feature set shipped through v1.7.3, including multi-agent meeting simulation, go/no-go investment verdicts with evidence decomposition, and customer-facing press release generation. No backend, no infrastructure. Remaining planned work (<code>/prfaq:import</code> , comparative analysis) is straightforward. Refinement risks limited to prompt tuning across Claude model updates and maintaining persona fidelity in the meeting and vote commands.
Viability	Medium	Two concerns, one survivable and one not. The survivable concern is maintainer capacity: at 500+ stars, estimated 4–8 hrs/week on issues, community, and prompt maintenance. The plugin has no servers, no API keys, no databases — just markdown files, a LaTeX template, and a shell script. If the maintainer disappears, the plugin keeps working and the codebase is trivially forkable. The non-survivable concern is platform dependency: the plugin runs exclusively on Claude Code, and Anthropic controls the plugin system, distribution, and competitive response (see FAQ 15). If Anthropic deprecates the plugin API, restricts third-party plugins, or ships a native alternative, the distribution channel is lost. The LaTeX output is portable — users retain their documents regardless — but the plugin’s ability to reach new users depends entirely on a platform it does not control. Mitigation: exploration of alternative CLI-based distribution channels (Feature 22) — not yet started, and filed as lower priority in the Feature Appendix pending evidence that Claude Code’s plugin ecosystem is becoming restrictive — plus an architecture that already minimizes platform-specific coupling, and the recognition that a positive Anthropic response (featuring the plugin) carries upside, while a native competing feature would validate the category even as it removes

Feature Appendix

Shipped (V1)

- F1. Interactive discovery workflow** — structured questions that guide the user from idea to document. Shipped v0.1.0.
- F2. Complete PR/FAQ document generation** — press release, external FAQs, internal FAQs, four risks assessment. Shipped v0.1.0.
- F3. LaTeX compilation to professional PDF** — shareable artifact, not a disposable brainstorm. Shipped v0.1.0.
- F4. Reference guides** — encoded best practices for PR structure, FAQ structure, four risks, common mistakes, unit economics, feasibility, and usability. Shipped v0.1.0.
- F5. Peer review agent (/prfaq:review)** — automated critique against Working Backwards principles and cognitive bias checklist. Shipped v0.2.0.
- F6. Research sub-agent (/prfaq:research)** — scans local ./research/ directory, indexed documents, and the web for evidence. Shipped v0.3.0.
- F7. Biblatex citation system** — academic-style citations for claims, linking evidence to assertions. Shipped v0.3.0.
- F8. /prfaq:feedback** — directed iteration from pointed user feedback; traces cascading effects via section dependency graph and surgically redrafts affected sections. Shipped v0.4.0.
- F9. /prfaq:feedback-to-us** — anonymous satisfaction feedback — 1–5 rating with optional comment to help improve the plugin. No document content transmitted. Shipped v0.7.0.
- F10. /prfaq:meeting** — simulates an Amazon-style PR/FAQ review meeting with four agentic personas (target customer, principal engineer, unit economics reviewer, UX bar raiser) plus a skeptical executive. Agent debate is visible to the user. Main Claude facilitates and synthesizes debate into pointed decision questions. The user is the PM and final decision-maker. Decisions feed into /prfaq:feedback for automated iteration. Shipped v0.6.0.
- F11. /prfaq:meeting-hive** — orchestrates the review meeting as autonomous multi-agent debate via Claude Code Agent Teams. Personas reach consensus without the user moderating every exchange, surfacing conflicts and agreements independently. Shipped v0.8.0; migrated to Agent Teams v1.2.0.
- F12. Stage awareness** — document stage (hypothesis, validated, growth) calibrates evidence expectations across all commands. Shipped v0.5.0.
- F13. Version tracking** — major/minor versioning with automatic increment after feedback cycles. Shipped v0.5.0.
- F14. Cross-references** — clickable \faqref and \featureref links between the press release, FAQs, and Feature Appendix. Shipped v0.4.0.
- F15. /prfaq:streamline** — scalpel editor that removes redundancy across sections, eliminates weasel words and hollow adjectives, compresses inflated phrases, and applies the “so what” test. Targets 10–20% length reduction without touching evidence, citations, or risk assessments. Shipped v0.8.1.
- F16. /prfaq:vote** — structured thinking exercise that walks through three questions: is the customer problem worth solving, is the solution differentiated, and is now the right time? Surfaces gaps in evidence, identifies cheap ways to reduce uncertainty, and renders a GO/NO-GO verdict. Not a mechanical formula — context always matters — but a discipline that forces the builder to confront what they know versus what they assume. Calibrations sourced from running vote on a real investment decision. Shipped v1.0.0; evidence decomposition and investment proportionality shipped v1.1.0.

- F17. `/prfaq:externalize`** — generates a customer-facing press release from the internal PR/FAQ and CHANGELOG. Detects release type (first release, major update, minor/patch), scopes content to what actually shipped, flags customer quotes for replacement with real testimonials. Version-stamped output. Shipped v1.0.0.
- F18. `/prfaq:export`** — Word document (.docx) output via pandoc as a lightweight alternative to PDF. Pre-processes custom LaTeX environments, resolves cross-references, styles output with Palatino serif and SectionBlue headings, centers title/subtitle, adds header/footer with stage and version. Requires only pandoc (~50 MB) instead of a full TeX distribution (~4 GB). Shipped v1.5.0.
- F19. `/prfaq:import`** — ingest an existing PR/FAQ draft (plain text, Markdown, or raw .tex), parse and structure it against the framework, research evidence for unsupported claims, format into the LaTeX template with proper environments and citations, compile to PDF, present a gap analysis. Shipped v1.0.0.

Planned

- F20. Comparative analysis mode** — evaluate two competing product ideas side-by-side
- F21. Template customization** — allow users to override the LaTeX template for branding
- F22. Alternative CLI-based distribution channels** — explore Codex CLI, Gemini CLI, and other terminal-based AI coding tools as additional distribution surfaces. The plugin's architecture — markdown skill files, a LaTeX template, and a shell script — is not deeply coupled to Claude Code's plugin system. Porting to other CLI-based agents would reduce single-platform concentration risk and expand reach to builders who use different tools. Not about supporting multiple LLM backends simultaneously; about ensuring the product-thinking workflow is available wherever builders work in terminals. Priority increases if Claude Code's plugin ecosystem becomes restrictive or if a competing CLI tool gains significant share among the target audience.

Won't Do

- F23. Web dashboard or SaaS platform** — the plugin runs locally inside Claude Code. Adding a web layer would introduce infrastructure, authentication, and hosting costs that contradict zero-ops design. If users want to share documents, they share the PDF.
- F24. Collaboration or multi-user editing** — PR/FAQ is a single-author process by design. The document is written by one person, then reviewed by others. Real-time collaboration would encourage design-by-committee, which Working Backwards explicitly avoids.
- F25. Project management features** — no task tracking, no roadmap views, no sprint planning. The plugin produces a decision document, not a project plan. Users who need project management have Linear, Jira, or beads.

References

- [1] Stack Overflow. *2025 Developer Survey*. Annual survey of 65,000+ developers on tools, practices, and AI adoption. 2025. URL: <https://survey.stackoverflow.co/2025>.
- [2] Andrej Karpathy. *Vibe Coding*. Original post coining the term “vibe coding” — viewed over 4.5 million times. Feb. 2025. URL: <https://x.com/karpathy/status/1886192184808149383>.
- [3] HFS Research and Publicis Sapient. *Smash Through Tech Debt: Why AI Is the Jackhammer*. Survey of 608 IT and business leaders estimating Global 2000 accumulated technical debt at \$1.5–2 trillion. 2025. URL: <https://www.hfsresearch.com/research/publicis-sapient-jackhammer/>.
- [4] Natallie Rocha. “This A.I. Tool Is Going Viral. Five Ways People Are Using It.” In: *The New York Times* (Jan. 2026). Profiles five non-programmers using Claude Code: school administrator, photographer, prosecutor, finance professor, welding business owner. URL: <https://www.nytimes.com/2026/01/23/technology/claude-code.html>.
- [5] Constellation Research. *Anthropic’s Claude Code Revenue Doubled Since Jan. 1*. Claude Code annualized run-rate revenue over \$2.5B as of February 2026, doubling since January 1. Anthropic total ARR is \$14B. Figure is Claude Code-specific. Feb. 2026. URL: <https://www.constellationr.com/insights/news/anthropics-claude-code-revenue-doubled-jan-1>.
- [6] TechCrunch. *Lovable Says It’s Nearing 8 Million Users*. Lovable, an AI coding startup founded in 2024, nearing 8 million users within one year. Nov. 2025. URL: <https://techcrunch.com/2025/11/10/lovable-says-its-nearing-8-million-users-as-the-year-old-ai-coding-startup-eyes-more-corporate-employees/>.
- [7] Amjad Masad. *Replit CEO: We Don’t Care About Professional Coders Anymore*. Replit’s 35M+ users; CEO explicitly pivoting away from professional developers toward non-coders. Jan. 2025. URL: <https://www.analyticsvidhya.com/blog/2025/01/replit-ceo-we-dont-care-about-professional-coders-anymore/>.
- [8] Sacra. *Bolt.new Revenue, Funding & News*. Bolt.new hit \$40M ARR in March 2025, progressing from \$4M ARR within 4 weeks of launch. 2025. URL: <https://sacra.com/c/bolt-new/>.
- [9] Sacra. *Cursor Revenue, Valuation & Funding*. Cursor reached \$1B ARR and 1M daily active users. 2025. URL: <https://sacra.com/c/cursor/>.
- [10] TechCrunch. *A Quarter of Startups in YC’s Current Cohort Have Codebases Almost Entirely AI-Generated*. 25% of Y Combinator W25 cohort startups have codebases that are almost entirely AI-generated. Mar. 2025. URL: <https://techcrunch.com/2025/03/06/a-quarter-of-startups-in-ycs-current-cohort-have-codebases-that-are-almost-entirely-ai-generated/>.
- [11] Colin Bryar and Bill Carr. *Working Backwards: Insights, Stories, and Secrets from Inside Amazon*. St. Martin’s Press, 2021. ISBN: 978-1250267597.
- [12] Marty Cagan. *Inspired: How to Create Tech Products Customers Love*. 2nd ed. Wiley, 2018. ISBN: 978-1119387503.
- [13] punt-labs. “Adoption signals — 2026-08-29”. Internal first-hand research note recording two direct product-manager conversations and inbound LinkedIn interest; see research/research-2026-08-29-adoption-signals.md in the prfaq repository. 2026.

- [14] Anthropic. *knowledge-work-plugins*. Official Cowork/Claude Code plugin marketplace launched January 2026 with 11 plugins, expanded to 24+ by February 2026; includes a product-management plugin for PRDs, roadmaps, and stakeholder communication. 2026. URL: <https://github.com/anthropics/knowledge-work-plugins>.
- [15] pmprompt. *claude-plugin-product-management*. Open-source Claude Code plugin bundling 26+ PM frameworks, including a /pmprompt:working-backwards command producing a PR/FAQ-style spec. 2026. URL: <https://github.com/pmprompt/claude-plugin-product-management>.
- [16] SemiAnalysis. *Claude Code is the Inflection Point*. 4% of GitHub public commits authored by Claude Code; projection of 20%+ by end of 2026. Feb. 2026. URL: <https://newsletter.semianalysis.com/p/claude-code-is-the-inflection-point>.